

Capstone 2: Biodiversity Project

Introduction

You are a biodiversity analyst working for the National Parks Service. You're going to help them analyze some data about species at various national parks.

Note: The data that you'll be working with for this project is *inspired* by real data, but is mostly fictional.

Step 1

Import the modules that you'll be using in this assignment:

- `from matplotlib import pyplot as plt`
- `import pandas as pd`

```
In [2]: from matplotlib import pyplot as plt
import pandas as pd
```

Step 2

You have been given two CSV files. `species_info.csv` with data about different species in our National Parks, including:

- The scientific name of each species
- The common names of each species
- The species conservation status

Load the dataset and inspect it:

- Load `species_info.csv` into a DataFrame called `species`

```
In [3]: species = pd.read_csv('species_info.csv')
```

Inspect each DataFrame using `.head()`.

```
In [4]: print species.head()
```

	category	scientific_name	common_names	conservation_status
0	Mammal	Clethrionomys gapperi gapperi	Gapper's Red-Backed Vole	NaN
1	Mammal	Bos bison	American Bison, Bison	NaN
2	Mammal	Bos taurus	Aurochs, Aurochs, Domestic Cattle (Feral), Dom...	NaN
3	Mammal	Ovis aries	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	NaN
4	Mammal	Cervus elaphus	Wapiti Or Elk	NaN

Step 3

Let's start by learning a bit more about our data. Answer each of the following questions.

How many different species are in the `species` DataFrame?

```
In [5]: species.scientific_name.nunique()
```

```
Out[5]: 5541
```

What are the different values of `category` in `species` ?

```
In [6]: species.category.unique()
```

```
Out[6]: array(['Mammal', 'Bird', 'Reptile', 'Amphibian', 'Fish', 'Vascular Plant',  
              'Nonvascular Plant'], dtype=object)
```

What are the different values of `conservation_status` ?

```
In [7]: species.conservation_status.unique()
```

```
Out[7]: array([nan, 'Species of Concern', 'Endangered', 'Threatened',  
              'In Recovery'], dtype=object)
```

Step 4

Let's start doing some analysis!

The column `conservation_status` has several possible values:

- `Species of Concern` : declining or appear to be in need of conservation
- `Threatened` : vulnerable to endangerment in the near future
- `Endangered` : seriously at risk of extinction
- `In Recovery` : formerly `Endangered` , but currently neither in danger of extinction throughout all or a significant portion of its range

We'd like to count up how many species meet each of these criteria. Use `groupby` to count how many `scientific_name` meet each of these criteria.

```
In [8]: species.groupby('conservation_status').scientific_name.nunique().reset_index()
```

Out[8]:

	conservation_status	scientific_name
0	Endangered	15
1	In Recovery	4
2	Species of Concern	151
3	Threatened	10

As we saw before, there are far more than 200 species in the `species` table. Clearly, only a small number of them are categorized as needing some sort of protection. The rest have `conservation_status` equal to `None`. Because `groupby` does not include `None`, we will need to fill in the null values. We can do this using `.fillna`. We pass in however we want to fill in our `None` values as an argument.

Paste the following code and run it to see replace `None` with `No Intervention`:

```
species.fillna('No Intervention', inplace=True)
```

```
In [9]: species.fillna('No Intervention', inplace=True)
```

Great! Now run the same `groupby` as before to see how many species require `No Protection`.

```
In [10]: species.groupby('conservation_status').scientific_name.nunique().reset_index()
```

Out[10]:

	conservation_status	scientific_name
0	Endangered	15
1	In Recovery	4
2	No Intervention	5363
3	Species of Concern	151
4	Threatened	10

Let's use `plt.bar` to create a bar chart. First, let's sort the columns by how many species are in each categories. We can do this using `.sort_values`. We use the keyword `by` to indicate which column we want to sort by.

Paste the following code and run it to create a new DataFrame called `protection_counts`, which is sorted by `scientific_name`:

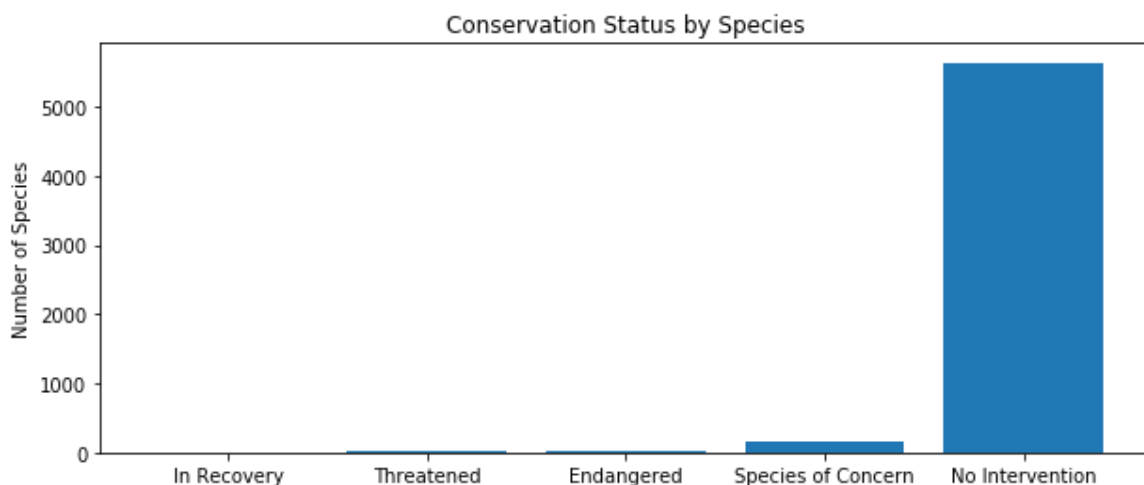
```
protection_counts = species.groupby('conservation_status')\
    .scientific_name.count().reset_index()\
    .sort_values(by='scientific_name')
```

```
In [11]: protection_counts = species.groupby('conservation_status')\
        .scientific_name.count().reset_index()\
        .sort_values(by='scientific_name')
```

Now let's create a bar chart!

1. Start by creating a wide figure with `figsize=(10, 4)`
2. Start by creating an axes object called `ax` using `plt.subplot()`.
3. Create a bar chart whose heights are equal to `scientific_name` column of `protection_counts`.
4. Create an x-tick for each of the bars.
5. Label each x-tick with the label from `conservation_status` in `protection_counts`
6. Label the y-axis Number of Species
7. Title the graph Conservation Status by Species
8. Plot the graph using `plt.show()`

```
In [12]: plt.figure(figsize=(10,4))
ax = plt.subplot()
plt.bar(range(len(protection_counts)),protection_counts.scientific_name)
ax.set_xticks(range(len(protection_counts)))
ax.set_xticklabels(protection_counts.conservation_status)
plt.ylabel('Number of Species')
plt.title('Conservation Status by Species')
plt.show()
```



Step 4

Are certain types of species more likely to be endangered?

Let's create a new column in `species` called `is_protected`, which is `True` if `conservation_status` is not equal to `No Intervention`, and `False` otherwise.

```
In [13]: species['is_protected'] = species.conservation_status.apply(lambda x: 'True' if x
```

Let's group by *both* `category` and `is_protected`. Save your results to `category_counts`.

```
In [14]: category_counts = species.groupby(['category', 'is_protected']).scientific_name.c
```

Examine category_count using head() .

```
In [15]: print category_counts.head()
```

	category	is_protected	scientific_name
0	Amphibian	False	73
1	Amphibian	True	7
2	Bird	False	442
3	Bird	True	79
4	Fish	False	116

It's going to be easier to view this data if we pivot it. Using pivot , rearrange category_counts so that:

- columns is conservation_status
- index is category
- values is scientific_name

Save your pivoted data to category_pivot . Remember to reset_index() at the end.

```
In [16]: category_pivot = category_counts.pivot(index='category', columns='is_protected',
```

Examine category_pivot .

```
In [17]: print category_pivot
```

is_protected	category	False	True
0	Amphibian	73	7
1	Bird	442	79
2	Fish	116	11
3	Mammal	176	38
4	Nonvascular Plant	328	5
5	Reptile	74	5
6	Vascular Plant	4424	46

Use the .columns property to rename the categories True and False to something more description:

- Leave category as category
- Rename False to not_protected
- Rename True to protected

```
In [18]: category_pivot = category_pivot.rename(columns={'False': 'not_protected', 'True':
```

Let's create a new column of category_pivot called percent_protected , which is equal to protected (the number of species that are protected) divided by protected plus not_protected (the total number of species).

```
In [19]: category_pivot['percent_protected'] = category_pivot.protected / (category_pivot.
```

Examine `category_pivot`.

```
In [20]: print category_pivot
```

	is_protected	category	not_protected	protected	percent_protected
0		Amphibian	73	7	0.087500
1		Bird	442	79	0.151631
2		Fish	116	11	0.086614
3		Mammal	176	38	0.177570
4		Nonvascular Plant	328	5	0.015015
5		Reptile	74	5	0.063291
6		Vascular Plant	4424	46	0.010291

It looks like species in category `Mammal` are more likely to be endangered than species in `Bird`. We're going to do a significance test to see if this statement is true. Before you do the significance test, consider the following questions:

- Is the data numerical or categorical?
- How many pieces of data are you comparing?

Based on those answers, you should choose to do a *chi squared test*. In order to run a chi squared test, we'll need to create a contingency table. Our contingency table should look like this:

	protected	not protected
Mammal	?	?
Bird	?	?

Create a table called `contingency` and fill it in with the correct numbers

```
In [21]: contingency = [[38, 176], [79, 442]]
```

In order to perform our chi square test, we'll need to import the correct function from `scipy`. Past the following code and run it:

```
from scipy.stats import chi2_contingency
```

```
In [22]: from scipy.stats import chi2_contingency
```

Now run `chi2_contingency` with `contingency`.

```
In [23]: chi2, pval, dof, expected = chi2_contingency(contingency)
print pval
```

```
0.445901703047197
```

It looks like this difference isn't significant!

Let's test another. Is the difference between `Reptile` and `Mammal` significant?

```
In [24]: contingency_2 = [[5, 74], [38, 176]]
chi2, pval, dof, expected = chi2_contingency(contingency_2)
print pval

0.02338465214871547
```

Yes! It looks like there is a significant difference between Reptile and Mammal!

Step 5

Conservationists have been recording sightings of different species at several national parks for the past 7 days. They've saved sent you their observations in a file called `observations.csv`. Load `observations.csv` into a variable called `observations`, then use `head` to view the data.

```
In [25]: import pandas as pd
observations = pd.read_csv('observations.csv')
print observations.head()
```

	scientific_name	park_name	observations
0	Vicia benghalensis	Great Smoky Mountains National Park	68
1	Neovison vison	Great Smoky Mountains National Park	77
2	Prunus subcordata	Yosemite National Park	138
3	Abutilon theophrasti	Bryce National Park	84
4	Githopsis specularioides	Great Smoky Mountains National Park	85

Some scientists are studying the number of sheep sightings at different national parks. There are several different scientific names for different types of sheep. We'd like to know which rows of `species` are referring to sheep. Notice that the following code will tell us whether or not a word occurs in a string:

```
In [26]: # Does "Sheep" occur in this string?
str1 = 'This string contains Sheep'
'Sheep' in str1
```

Out[26]: True

```
In [27]: # Does "Sheep" occur in this string?
str2 = 'This string contains Cows'
'Sheep' in str2
```

Out[27]: False

Use `apply` and a `lambda` function to create a new column in `species` called `is_sheep` which is True if the `common_names` contains 'Sheep', and False otherwise.

```
In [28]: species['is_sheep'] = species.common_names.apply(lambda x: 'Sheep' in x)
```

Select the rows of `species` where `is_sheep` is True and examine the results.

```
In [29]: species[species.is_sheep]
```

Out[29]:

	category	scientific_name	common_names	conservation_status	is_protected	is_sheep
3	Mammal	Ovis aries	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	False	True
1139	Vascular Plant	Rumex acetosella	Sheep Sorrel, Sheep Sorrell	No Intervention	False	True
2233	Vascular Plant	Festuca filiformis	Fineleaf Sheep Fescue	No Intervention	False	True
3014	Mammal	Ovis canadensis	Bighorn Sheep, Bighorn Sheep	Species of Concern	True	True
3758	Vascular Plant	Rumex acetosella	Common Sheep Sorrel, Field Sorrel, Red Sorrel,...	No Intervention	False	True
3761	Vascular Plant	Rumex paucifolius	Alpine Sheep Sorrel, Fewleaved Dock, Meadow Dock	No Intervention	False	True
4091	Vascular Plant	Carex illota	Sheep Sedge, Smallhead Sedge	No Intervention	False	True
4383	Vascular Plant	Potentilla ovina var. ovina	Sheep Cinquefoil	No Intervention	False	True
4446	Mammal	Ovis canadensis sierrae	Sierra Nevada Bighorn Sheep	Endangered	True	True

Many of the results are actually plants. Select the rows of `species` where `is_sheep` is `True` and `category` is `Mammal`. Save the results to the variable `sheep_species`.

```
In [30]: sheep_species = species[(species.is_sheep) & (species.category == 'Mammal')]
print sheep_species
```

	category	scientific_name	common_names	conservation_status	is_protected	is_sheep
3	Mammal	Ovis aries	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	False	True
3014	Mammal	Ovis canadensis	Bighorn Sheep, Bighorn Sheep	Species of Concern	True	True
4446	Mammal	Ovis canadensis sierrae	Sierra Nevada Bighorn Sheep	Endangered	True	True

Now merge `sheep_species` with `observations` to get a `DataFrame` with observations of sheep. Save this `DataFrame` as `sheep_observations`.


```
In [31]: sheep_observations = pd.merge(sheep_species, observations)
print sheep_observations
```

	category	scientific_name	\
0	Mammal	Ovis aries	
1	Mammal	Ovis aries	
2	Mammal	Ovis aries	
3	Mammal	Ovis aries	
4	Mammal	Ovis canadensis	
5	Mammal	Ovis canadensis	
6	Mammal	Ovis canadensis	
7	Mammal	Ovis canadensis	
8	Mammal	Ovis canadensis sierrae	
9	Mammal	Ovis canadensis sierrae	
10	Mammal	Ovis canadensis sierrae	
11	Mammal	Ovis canadensis sierrae	

	common_names	conservation_status	\
0	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	
1	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	
2	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	
3	Domestic Sheep, Mouflon, Red Sheep, Sheep (Feral)	No Intervention	
4	Bighorn Sheep, Bighorn Sheep	Species of Concern	
5	Bighorn Sheep, Bighorn Sheep	Species of Concern	
6	Bighorn Sheep, Bighorn Sheep	Species of Concern	
7	Bighorn Sheep, Bighorn Sheep	Species of Concern	
8	Sierra Nevada Bighorn Sheep	Endangered	
9	Sierra Nevada Bighorn Sheep	Endangered	
10	Sierra Nevada Bighorn Sheep	Endangered	
11	Sierra Nevada Bighorn Sheep	Endangered	

	is_protected	is_sheep	park_name	observations
0	False	True	Yosemite National Park	126
1	False	True	Great Smoky Mountains National Park	76
2	False	True	Bryce National Park	119
3	False	True	Yellowstone National Park	221
4	True	True	Yellowstone National Park	219
5	True	True	Bryce National Park	109
6	True	True	Yosemite National Park	117
7	True	True	Great Smoky Mountains National Park	48
8	True	True	Yellowstone National Park	67
9	True	True	Yosemite National Park	39
10	True	True	Bryce National Park	22
11	True	True	Great Smoky Mountains National Park	25

How many total sheep observations (across all three species) were made at each national park? Use `groupby` to get the `sum` of `observations` for each `park_name`. Save your answer to `obs_by_park`.

This is the total number of sheep observed in each park over the past 7 days.

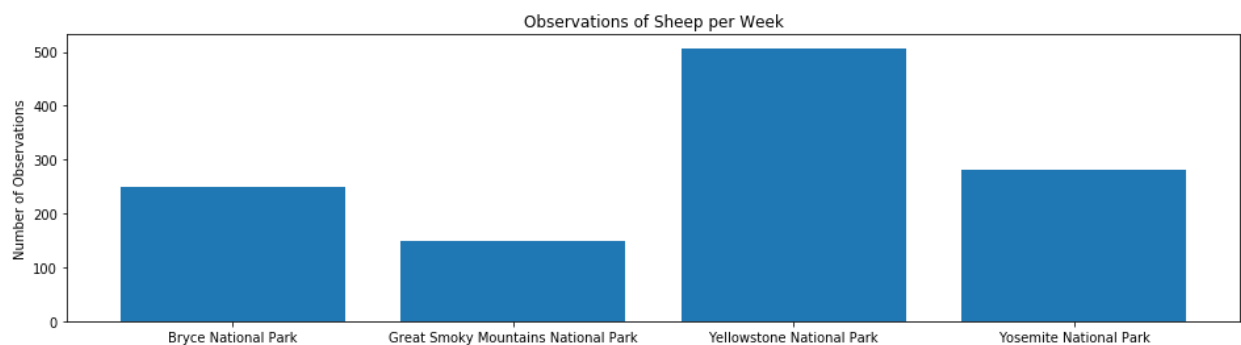
```
In [32]: obs_by_park = sheep_observations.groupby('park_name').observations.sum().reset_index
print obs_by_park
```

	park_name	observations
0	Bryce National Park	250
1	Great Smoky Mountains National Park	149
2	Yellowstone National Park	507
3	Yosemite National Park	282

Create a bar chart showing the different number of observations per week at each park.

1. Start by creating a wide figure with `figsize=(16, 4)`
2. Start by creating an axes object called `ax` using `plt.subplot()`.
3. Create a bar chart whose heights are equal to `observations` column of `obs_by_park`.
4. Create an x-tick for each of the bars.
5. Label each x-tick with the label from `park_name` in `obs_by_park`
6. Label the y-axis Number of Observations
7. Title the graph Observations of Sheep per Week
8. Plot the graph using `plt.show()`

```
In [33]: plt.figure(figsize=(16, 4))
ax=plt.subplot()
plt.bar(range(len(obs_by_park)),obs_by_park.observations)
ax.set_xticks(range(len(obs_by_park)))
ax.set_xticklabels(obs_by_park.park_name)
plt.ylabel('Number of Observations')
plt.title('Observations of Sheep per Week')
plt.show()
```



Our scientists know that 15% of sheep at Bryce National Park have foot and mouth disease. Park rangers at Yellowstone National Park have been running a program to reduce the rate of foot and mouth disease at that park. The scientists want to test whether or not this program is working. They want to be able to detect reductions of at least 5 percentage point. For instance, if 10% of sheep in Yellowstone have foot and mouth disease, they'd like to be able to know this, with confidence.

Use the sample size calculator at [Optimizely \(https://www.optimizely.com/sample-size-calculator/\)](https://www.optimizely.com/sample-size-calculator/) to calculate the number of sheep that they would need to observe from each park. Use the default level of significance (90%).

Remember that "Minimum Detectable Effect" is a percent of the baseline.

```
In [34]: '''Baseline = 15
Minimum_detectable_effect = 100*5 percentage points/baseline = 33.33%
Statistical significance = 90%
Sample_Size_Per_Park = 510
## From sample size calculator at Optimizely.'''
```

```
Out[34]: 'Baseline = 15\nMinimum_detectable_effect = 100*5 percentage points/baseline =
33.33%\nStatistical significance = 90%\nSample_Size_Per_Park = 510\n## From sa
mple size calculator at Optimizely.'
```

How many weeks would you need to observe sheep at Bryce National Park in order to observe enough sheep? How many weeks would you need to observe at Yellowstone National Park to observe enough sheep?

```
In [35]: '''Bryce Weeks Observing = sample_size_per_park/Bryce observations
510/250 = 2.04 weeks
We need to spend 2.04 weeks at Bryce National Park in order to observe enough she
'''Yellowstone Weeks Observing = sample_size_per_park/Yellowstone observations
510/507 = 1.00 week
We need to spend 1 week at Yellowstone National Park in order to observe enough s
```

```
Out[35]: 'Yellowstone Weeks Observing = sample_size_per_park/Yellowstone observations\n
510/507 = 1.00 week\nWe need to spend 1 week at Yellowstone National Park in o
rder to observe enough sheep.'
```

```
In [ ]:
```